

Homework 5: Optical Flow & Tracking

CS 566, Fall 2025

Total points: 7

Due on Tuesday, December 2, 2025

Please follow the homework guidelines at <https://pages.cs.wisc.edu/~mohitg/courses/CS566/hw-guidelines.html> in your submission.

`runHw5.py` will be your main interface for running your code. Parameters for the different programs or unit tests can also be set in that file. Before submission, make sure that you can run all your programs with the command `python runHw5.py` all with no errors.

This assignment makes use of the following Python packages: NumPy, Matplotlib, scikit-image, scikit-learn, scipy.

Challenge 1: Optical Flow (6 points)

In this challenge you are asked to develop an optical flow system. You are given a sequence of 6 images (`flow1.png-flow6.png`) of a dynamic scene. Your task is to develop an algorithm that computes optical flow estimates at each image point using the 5 pairs (1 & 2, 2 & 3, 3 & 4, 4 & 5, 5 & 6) of consecutive images.

Optical flow estimates can be computed using the optical flow constraint equation and Lucas-Kanade solution presented in class. For smooth motions, this algorithm should produce robust flow estimates. However, given that the six images were taken with fairly large time intervals in between consecutive images, the brightness and temporal derivatives used by the algorithm are expected to be unreliable. Therefore, you are advised to implement a different (and simpler) optical flow algorithm.

Given two consecutive images (say 1 and 2), establish correspondence between points in the two images using template matching. For each image point in the first image, take a small window (say 7×7) around the point and use it as the template to find the same point in the second image. While searching for the corresponding point in the second image, you can confine the search to a small window around the pixel in the second image that has the same coordinates as the one in the first image. The center of the 7×7 image window in the second image that is maximally correlated with the 7×7 window in the first image is assumed to be the corresponding point. The vector between two corresponding points is the optical flow (u,v). To compute the correlation, we recommend using the function `skimage.feature.match_template`.

Compute Flow (5 points)

Write a program `computeFlow` that computes optical flow between two gray-level images, and produces the optical flow vector field as a "needle map" of a given resolution, overlaid on the first

of the two images.

```
result = computeFlow(img1, img2, search_half_window_size, ...
                     template_half_window_size, grid_MN)
```

We have already written the skeleton of the `computeFlow` function for you. In the `computeFlow` function you will see a double for loop. Within that for loop you will have to complete the remainder parts of the code that:

1. Extract the template and `search_window` from `img1` and `img2`.
2. Find the location in the `search_window` that best matches the template. Recall that normalized cross-correlation is what we use in template matching to measure how well the template matches.
3. Given the location of the best match, store the optical flow (motion in `x` and `y`).

You may use the `scikit-image` function `match_template` to perform the normalized cross-correlation. The needle map (and hence optical flow) should only be computed and drawn on an $M \times N$ grid equally sampled over the image.

At the end of the `computeFlow` function we have added code that overlays the computed optical flow as a needle map on top of the image. We use the `matplotlib` function `quiver` to annotate the image with the computed optical flow. You need to choose a value for the grid spacing that gives good results without taking excessively long to compute.

Synthetic case (1 point)

For debugging purposes use the test case in `debugla`. In this synthetic case, the flow field consists of horizontal vectors of the same magnitude (translational motion parallel to the image plane). Note that in the real case, foreshortening effects, occlusions, and reflectance variations (as well as noise) complicate the result. If you have implemented the `computeFlow` function correctly the `debugla` code should output an image that looks like this:




Figure 1: Example output for the synthetic debug case.

Challenge 2: Tracking (8 points)

Your task is to develop a vision system that tracks the location of an object across video frames. Object tracking is a challenging problem since an object's appearance, pose and scale tend to change as time progresses. In class we have discussed three popular tracking methods: template-based tracking, histogram-based tracking and detection-based tracking. In this challenge, we will assume the color distribution of an object stays relatively constant over time. Therefore, we will track an object using its color histogram.

A color histogram describes the color distribution of a color image. The color histogram that you will need to compute is defined as follows. Each bin of the color histogram represents a range of colors, and the number of votes in each bin indicates the number of pixels that have the colors within the corresponding color range.

Use the provided function `rgb2ind` to divide the color range of an image into a number of bins. You can specify the number of bins as an input to `rgb2ind`. `rgb2ind` then returns a color map along with an index image, in which the color of each pixel is represented as an index to the color map. After computing an index image, you can use the NumPy function `histogram` to compute the color histogram of the image.

Be careful: In the initialization of each challenge, you should generate a color map from the object of interest once, and compute all subsequent color histograms based on **the same** color map. It is only meaningful to compare two histograms computed on the same color map.

Write a program named `trackingTester` that estimates the location of an object in video frames.

`trackingTester(data_params, tracking_params)`

`trackingTester` should draw a box around the target in each video frame, and save all the annotated video frames as PNGs into a sub-folder given in `data_params.out_dir`. Use `drawBox.m` (included in the homework package) to draw a box on images.

We have already written the skeleton of the `trackingTester` function for you. Please read the function and fill in the necessary code in the places indicated by the comments. In the `trackingTester` function we provide you will still need to complete the code that does the following:

1. Extract the rectangular region of interest from the first frame and create its histogram.
2. For each subsequent frame
 - (a) Extract the search window
 - (b) Create a histogram for all possible candidate regions in the search window. Hint: You may also find the function `skimage.util.view_as_windows` useful.
 - (c) Find the histogram that best matches the one you originally computed
 - (d) Update the location of the rectangular bounding box that tells us where the object we are tracking is at.

At the end of `challenge2a`, `challenge2b`, and `challenge2c`, an output video is created by calling the `generate_video` function. **IMPORTANT:** This video should be in the original color of the images (not the indexed color)! You only need the index conversion for generating the histograms, so you should do that with copies of the frames.

Include all the code you have written, as well as the generated video of the annotated frames in your submission. **DO NOT** include the three tracking datasets, nor the individual output frames, as they are quite large.